

Credit Card Fraud Detection

Namrata Kankaria

Abstract

- **Dataset** - Credit Card dataset is used in this project.
- **Question** - How can we detect fraudulent credit card transactions.
- **Methods used** - Pandas, Matplotlib, Machine Learning Models, Classification Report and Confusion Matrix.
- **Findings** - Number of fraud and valid transactions in the dataset, Amount involved in fraud transactions, Correlations of features in the dataset, Comparing different Machine Learning models.

Motivation

- Credit card fraud costs consumers and the financial company billions of dollars annually, and fraudsters continuously try to find new rules and tactics to commit illegal actions. Thus, fraud detection systems have become essential for banks and financial institution to minimize their losses.
- Through this project we aim to recognize fraudulent credit card transactions so that the customers of credit card companies are not charged for items that they did not purchase.

Dataset(s)

- Credit Card dataset is used in this project which contains transactions made via credit cards in September 2013 by European cardholders.
- This dataset presents 492 frauds out of 284,807 transactions.
- The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.
- I found this dataset from kaggle.

Data Preparation and Cleaning

To prepare the data for analysis, I did the following –

- Checked if there is any null or NA records. If yes, then dropped those records.
- Bifurcated Training and Testing Data to train the model and then evaluate the performance of machine learning models.

Research Question(s)

How to detect fraudulent credit card transactions so that the customers of credit card companies are not charged for items that they did not purchase?

- Identify number of fraud and valid transactions.
- Identify the details of amount involved in fraud transactions.
- Identify which features in the dataset correlate with each which can be helpful in choosing those features that are most relevant for the prediction.

Research Question(s)

- Compare different Machine Learning models and evaluate them.
- Is this a problem of high recall or high precision.

Methods

- Pandas is used to retrieve and manipulate data from the dataset.
- After data preparation and cleaning, Matplotlib is used for data visualization.
- Machine learning models – Logistic Regression, Decision Tree and Random Forest is used for detecting fraud credit card transactions.
- For training the Machine Learning model, data is split into two parts- Training data and Testing data.
- For evaluating classifier, Classification Report and Confusion Matrix is used.

Findings

We can see most of the features of the dataset are clustered around 0 with some or no outliers.



Findings


Only 0.17%
fraudulent
transaction out of
all the transactions.
The data is highly
unbalanced.

```
[→ Fraudulent Transactions: 492  
Valid Transactions: 284315  
Percentage of Fraudulent Transactions: 0.17%
```

Findings


50% of fraud transactions are of less than \$9.

Only 5% of fraud transactions are more than \$648.



Findings

From the HeatMap, we can clearly see that there are no strongly correlated features in the dataset. There are some features that either has a positive or a negative correlation with each other. Other than that, all the correlation values lie somewhere in the neutral or zero-range.



Findings

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56856
1	0.94	0.75	0.84	106
accuracy			1.00	56962
macro avg	0.97	0.88	0.92	56962
weighted avg	1.00	1.00	1.00	56962

Random Tree

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56856
1	0.87	0.73	0.79	106
accuracy			1.00	56962
macro avg	0.93	0.86	0.89	56962
weighted avg	1.00	1.00	1.00	56962


	precision	recall	f1-score	support
0	1.00	1.00	1.00	56856
1	0.83	0.58	0.69	106
accuracy			1.00	56962
macro avg	0.91	0.79	0.84	56962
weighted avg	1.00	1.00	1.00	56962

Logistic Regression

Classification Report

Decision Tree

Findings



Confusion Matrix

Limitations

The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions. The data is biased.

Conclusions

- Only 0.17% fraudulent transaction out of all the transactions. The data is biased.
- The features of the dataset are largely uncorrelated.
- 50% of fraud transactions are of less than \$9.
- Only 5% of fraud transactions are more than \$648.
- Random Forest Model is providing a better result even for the recall.
- This is a problem of high recall as there is a higher cost associated with false negative than false positive. If a fraudulent transaction is labelled as legitimate transaction, it will cause a huge loss to the company and their customers.

Acknowledgements

As the dataset is hosted on Kaggle, I would like to thank entire kaggle community for the knowledge sharing and resource.

References

- UC San DiegoX: DSE200x - Python for Data Science. Thanks to professors - Ilkay Altintas and Leo Porter.
- Kaggle <https://www.kaggle.com/mlg-ulb/creditcardfraud>

Importing all the necessary libraries

```
In [29]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import gridspec
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score
from sklearn.metrics import precision_score, recall_score
from sklearn.metrics import f1_score, matthews_corrcoef
from sklearn.metrics import confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

Loading the data

```
In [57]: data = pd.read_csv('./creditcard.csv', sep=',')
```

Exploring Data

```
In [58]: data.head(10)
```

Out[58]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	
5	2.0	-0.425966	0.960523	1.141109	-0.168252	0.420987	-0.029728	0.476201	0.260314	-
6	4.0	1.229658	0.141004	0.045371	1.202613	0.191881	0.272708	-0.005159	0.081213	
7	7.0	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631	-3.807864	
8	7.0	-0.894286	0.286157	-0.113192	-0.271526	2.669599	3.721818	0.370145	0.851084	-
9	9.0	-0.338262	1.119593	1.044367	-0.222187	0.499361	-0.246761	0.651583	0.069539	-

10 rows × 31 columns

In [59]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null  float64
1   V1       284807 non-null  float64
2   V2       284807 non-null  float64
3   V3       284807 non-null  float64
4   V4       284807 non-null  float64
5   V5       284807 non-null  float64
6   V6       284807 non-null  float64
7   V7       284807 non-null  float64
8   V8       284807 non-null  float64
9   V9       284807 non-null  float64
10  V10      284807 non-null  float64
11  V11      284807 non-null  float64
12  V12      284807 non-null  float64
13  V13      284807 non-null  float64
14  V14      284807 non-null  float64
15  V15      284807 non-null  float64
16  V16      284807 non-null  float64
17  V17      284807 non-null  float64
18  V18      284807 non-null  float64
19  V19      284807 non-null  float64
20  V20      284807 non-null  float64
21  V21      284807 non-null  float64
22  V22      284807 non-null  float64
23  V23      284807 non-null  float64
24  V24      284807 non-null  float64
25  V25      284807 non-null  float64
26  V26      284807 non-null  float64
27  V27      284807 non-null  float64
28  V28      284807 non-null  float64
29  Amount   284807 non-null  float64
30  Class    284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

From the above results we can observe that there are total 284807 rows/records and 31 columns

```
In [60]: data.describe()
```

Out[60]:

	Time	V1	V2	V3	V4	V5
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.919560e-15	5.688174e-16	-8.769071e-15	2.782312e-15	-1.552563e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01

8 rows × 31 columns

Data Cleaning

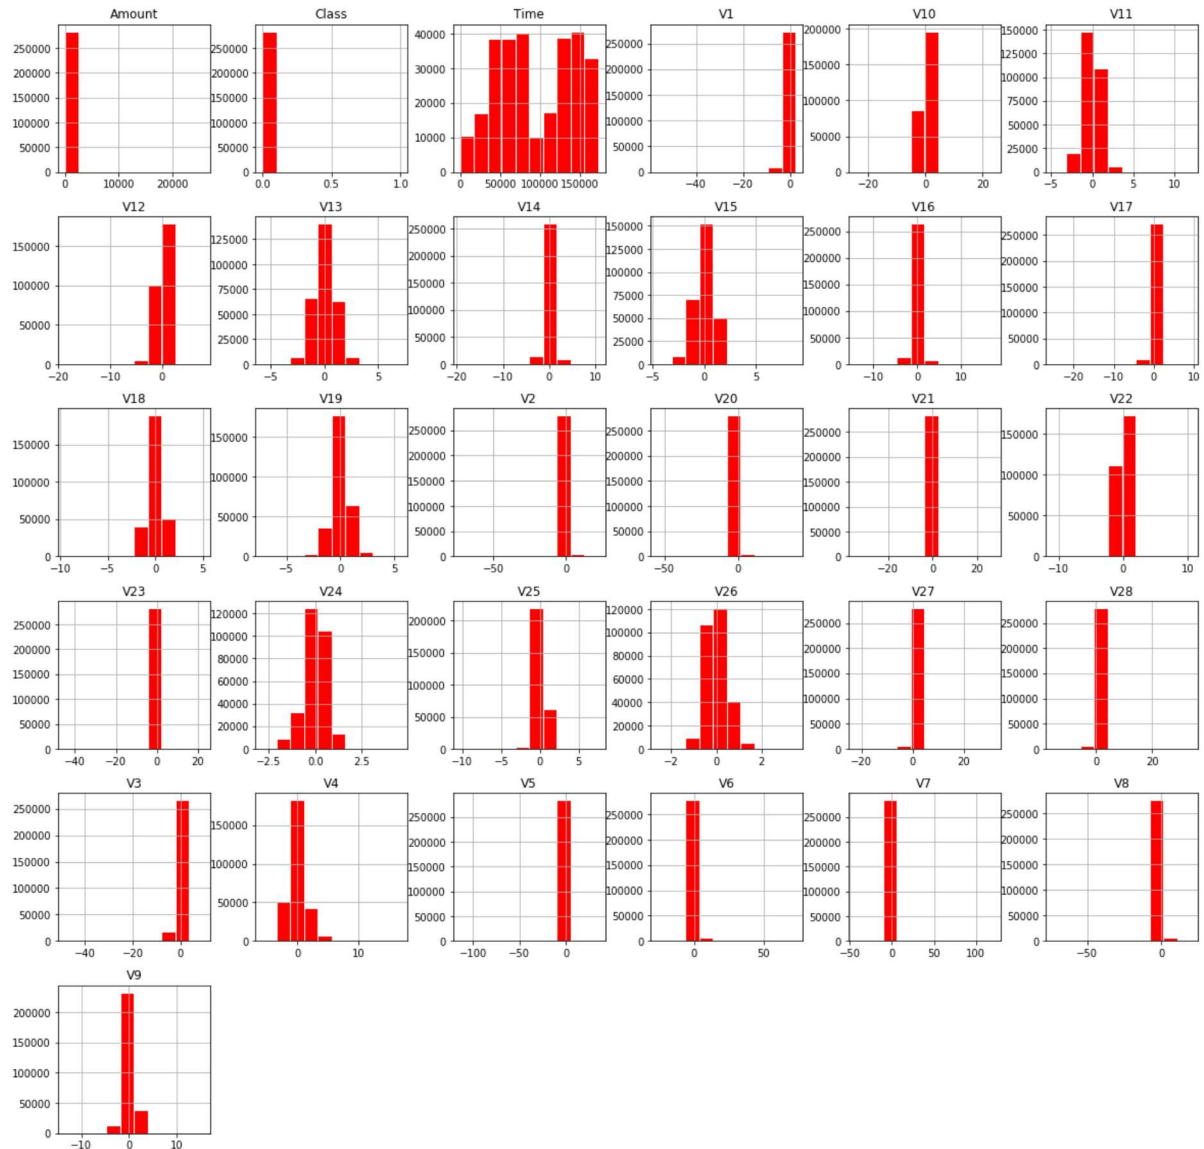
```
In [79]: data.isnull().sum()
```

```
Out[79]: Time      0  
V1      0  
V2      0  
V3      0  
V4      0  
V5      0  
V6      0  
V7      0  
V8      0  
V9      0  
V10     0  
V11     0  
V12     0  
V13     0  
V14     0  
V15     0  
V16     0  
V17     0  
V18     0  
V19     0  
V20     0  
V21     0  
V22     0  
V23     0  
V24     0  
V25     0  
V26     0  
V27     0  
V28     0  
Amount  0  
Class   0  
dtype: int64
```

No Null records

Plotting histogram of each feature

```
In [80]: data.hist(figsize = (20, 20),color = 'red' , edgecolor = 'white' )
plt.show()
```



We can see most of the feature V's are clustered around 0 with some or no outliers.

Identifying number of fraud and valid cases transactions

```
In [81]: fraud_cases = data[data['Class'] == 1]
valid_cases = data[data['Class'] == 0]
print('Fraudulent Transactions: {}'.format(len(fraud_cases)))
print('Valid Transactions: {}'.format(len(valid_cases)))
print('Percentage of Fraudulent Transactions: {:.2%}'.format((fraud_cases.shape[0] / data.shape[0])))
```

Fraudulent Transactions: 492

Valid Transactions: 284315

Percentage of Fraudulent Transactions: 0.17%

Inspecting more on Fraudulent Transactions and Valid Transactions

Amount Details of the Fraudulent Transaction

```
In [82]: fraud_cases.Amount.describe()
```

```
Out[82]: count      492.000000
mean       122.211321
std        256.683288
min         0.000000
25%         1.000000
50%         9.250000
75%        105.890000
max       2125.870000
Name: Amount, dtype: float64
```

Amount Details of the Valid Transaction

```
In [83]: valid_cases.Amount.describe()
```

```
Out[83]: count      284315.000000
mean         88.291022
std         250.105092
min          0.000000
25%          5.650000
50%         22.000000
75%         77.050000
max       25691.160000
Name: Amount, dtype: float64
```



```
In [84]: fraud_cases.Amount.tail(25)
```

```
Out[84]: 254395      7.59
          255403      4.97
          255556      0.77
          258403    296.00
          261056     45.51
          261473      4.90
          261925    156.00
          262560      4.69
          262826      0.77
          263080      1.00
          263274      0.77
          263324    127.14
          263877      0.38
          268375     39.98
          272521     12.31
          274382      0.00
          274475     39.90
          275992    634.30
          276071     19.95
          276864    349.08
          279863    390.00
          280143      0.76
          280149     77.89
          281144    245.00
          281674     42.53
          Name: Amount, dtype: float64
```

```
In [85]: fraud_cases.sort_values("Amount", axis = 0, ascending = True, inplace = True,
na_position = 'last')
percentage = [(x/491)*100 for x in range(492)]
fraud_amount = fraud_cases.Amount
plt.figure(figsize=(10,5))


plt.plot(fraud_amount, percentage,color='green')
plt.axvline(122, color='black',linestyle='--',linewidth=1,label='Mean= 122')
plt.legend()
plt.text(130,3, 'Mean')
plt.title('Number of Fraud Transactions Vs Amount')
plt.xlabel('Amount')
plt.ylabel('% Number of Transactions')
plt.grid()
plt.yticks(np.arange(0,110,10))
plt.ylim(0,105)
plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


"""Entry point for launching an IPython kernel.



Plotting the Correlation Matrix

The correlation matrix graphically gives us an idea of how features correlate with each other and can help us predict what are the features that are most relevant for the prediction.

```
In [86]: correlation_matrix = data.corr()
fig = plt.figure(figsize = (12, 9))
sns.heatmap(correlation_matrix, vmax = .7, square = True, cmap="YlGnBu")
plt.show()
```



From the HeatMap, we can clearly see that there are no strongly correlated features in the dataset. There are some features that either has a positive or a negative correlation with each other. Other than that, all the correlation values lie somewhere in the neutral or zero-range.

Applying Machine Learning Methods

Bifurcation of Training and Testing Data

We will be dividing the dataset into two main groups. One for training the model and the other for Testing our trained model's performance.

```
In [87]: X = data.drop(['Class'], axis = 1)
Y = data["Class"]
x_data = X.values
y_data = Y.values
xTrain, xTest, yTrain, yTest = train_test_split(x_data, y_data, test_size = 0.2, random_state = 43)
```

Random Forest Model

```
In [88]: random_forest = RandomForestClassifier()
random_forest.fit(xTrain, yTrain)
y_prediction = random_forest.predict(xTest)
```

Evaluating the classifier


```
In [89]: print(classification_report(yTest, y_prediction))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56856
1	0.93	0.75	0.83	106
accuracy			1.00	56962
macro avg	0.96	0.87	0.91	56962
weighted avg	1.00	1.00	1.00	56962

Visualizing the Confusion Matrix

```
In [90]: fig, ax = plt.subplots()
sns.heatmap(confusion_matrix(yTest, y_prediction, normalize='true'), annot=True, ax=ax)
ax.set_title("Confusion Matrix")
ax.set_ylabel("Real Value")
ax.set_xlabel("Predicted")

plt.show()
```



Logistic Regression

```
In [91]: regress = LogisticRegression(solver="lbfgs", max_iter=1000).fit(xTrain, yTrain)
y_prediction = regress.predict(xTest)
```

Evaluating the classifier


```
In [92]: print(classification_report(yTest, y_prediction))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56856
1	0.83	0.58	0.69	106
accuracy			1.00	56962
macro avg	0.91	0.79	0.84	56962
weighted avg	1.00	1.00	1.00	56962

Visualizing the Confusion Matrix

```
In [93]: fig, ax = plt.subplots()
sns.heatmap(confusion_matrix(yTest, y_prediction, normalize='true'), annot=True, ax=ax)
ax.set_title("Confusion Matrix")
ax.set_ylabel("Real Value")
ax.set_xlabel("Predicted")

plt.show()
```



Decision Tree

```
In [94]: decision_tree = DecisionTreeClassifier(max_depth=4, criterion="entropy")
decision_tree.fit(xTrain, yTrain)
y_prediction = decision_tree.predict(xTest)
```

Evaluating the classifier

```
In [95]: print(classification_report(yTest, y_prediction))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56856
1	0.87	0.73	0.79	106
accuracy			1.00	56962
macro avg	0.93	0.86	0.89	56962
weighted avg	1.00	1.00	1.00	56962

Visualizing the Confusion Matrix

```
In [96]: fig, ax = plt.subplots()
sns.heatmap(confusion_matrix(yTest, y_prediction, normalize='true'), annot=True, ax=ax)
ax.set_title("Confusion Matrix")
ax.set_ylabel("Real Value")
ax.set_xlabel("Predicted")

plt.show()
```

